

INVARIANCE AUDITS FOR ROBOT POLICIES: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study robot policy evaluation. The motivating bet is: Audit which policy invariances are physically valid and which erase necessary distinctions.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For robot policy evaluation, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2013; 2018; 2025; 2026; 2021; 2019). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.537	0.463
observed_only_wam	hidden_mode_shift	0.185	0.815
augmentation_as_data_wam	nominal	0.667	0.333
augmentation_as_data_wam	hidden_mode_shift	0.250	0.750
scalar_uncertainty_wam	nominal	0.652	0.348
scalar_uncertainty_wam	hidden_mode_shift	0.315	0.685
proposed_branch_mechanism	nominal	0.748	0.252
proposed_branch_mechanism	hidden_mode_shift	0.573	0.427

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Invariance Audits for Robot Policies argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Sqp-based mobile manipulator motion planning with controlled infeasibility for physically valid task failure, 2013. <https://doi.org/10.1115/detc2013-13377>.
- The design and evaluation of an ergonomic contactless gesture control system for industrial robots, 2018. <https://doi.org/10.1155/2018/9791286>.
- Evaluation of image enhancement techniques for vision-based navigation under low illumination, 2019. <https://doi.org/10.1155/2019/5015741>.
- Cobot user frame calibration: Evaluation and comparison between positioning repeatability performances achieved by traditional and vision-based methods, 2021. <https://doi.org/10.3390/robotics10010045>.
- Coordinated robustness evaluation framework for vision-language models, 2025. <http://arxiv.org/abs/2506.05429v1>.
- The evaluation bottleneck of vision-language-action models: A evaluation-centric survey, 2026. <https://doi.org/10.20944/preprints202606.0425.v1>.